
GNN-Based BERT for Understanding Context from Music

Tabarak Tasnim Tanu¹

Abstract

Music is a multi-layered signal containing complementary acoustic and semantic information. This project investigates a multimodal approach for music context understanding by representing audio tracks as graphs and combining graph-based acoustic representations with Transformer-based textual representations. Each track is divided into 5-second segments represented by 128-bin log-Mel features, with temporal and cosine-similarity edges forming the audio graph. A frozen DistilBERT encoder represents track titles, while GraphSAGE encodes the audio graph; their representations are integrated using both simple feature fusion and a cross-attention mechanism. Experiments on the FMA-small dataset using eight genre classes show that the Cross-Attention GNN+BERT model achieves 40.50% test accuracy, 41.18% Macro-F1, and 43.25% Macro AUC-PR, outperforming the individual DistilBERT, CNN, and GraphSAGE baselines. The results demonstrate that multimodal fusion provides the main performance improvement, while the additional benefit of cross-attention over simple fusion is comparatively modest.

1. Introduction

Music is a multimodal domain where meaningful context is expressed through both acoustic and textual information. Acoustic properties such as timbre, rhythm, harmony, and spectral patterns provide information about musical structure and genre, while titles, lyrics, tags, and captions provide complementary semantic cues.

Graph Neural Networks (GNNs) can represent relationships among structured audio segments. By connecting consecutive segments and acoustically similar segments, a graph allows information to be propagated across related parts of a track. In this work, GraphSAGE is used to learn track-level representations from these audio graphs (Hamilton et al., 2017).

Transformer-based language models provide contextual representations of text. BERT introduced bidirectional context-

tual representations (Devlin et al., 2019), while DistilBERT offers a smaller and more computationally efficient alternative (Sanh et al., 2019). Here, DistilBERT is used to encode track titles as complementary semantic features.

This project investigates whether combining graph-based audio and textual representations improves music genre classification. Five models are evaluated: DistilBERT-only, CNN, GraphSAGE, simple GNN-BERT fusion, and cross-attention GNN-BERT. Experiments are conducted on the FMA-small dataset (Defferrard et al., 2017) using eight genre classes. Evaluation includes accuracy, Macro-F1, Micro-F1, Macro AUC-PR, class-level analysis, t-SNE visualization, and case studies. The results demonstrate that multimodal fusion improves performance over unimodal models, with cross-attention providing a smaller additional gain.

2. Related Work

2.1. Music Information Retrieval

Music information retrieval (MIR) focuses on extracting meaningful information from audio recordings, including genre, mood, tags, and other semantic properties. The Free Music Archive (FMA) provides a large collection of openly available music together with metadata and standardized genre information, making it suitable for reproducible music classification experiments (Defferrard et al., 2017). In this project, the FMA-small subset is used to investigate genre classification from both acoustic and textual information.

2.2. Graph Neural Networks for Audio Representation

Graph Neural Networks provide a framework for learning representations from data with explicit relational structures. Unlike conventional neural networks that primarily operate on fixed grids or independent feature vectors, GNNs propagate information between connected nodes. GraphSAGE introduces an inductive neighborhood aggregation framework in which node representations are learned by aggregating information from local neighborhoods (Hamilton et al., 2017).

For music representation, graph structures can naturally capture relationships between different temporal segments of an audio track. Consecutive segments can be connected to

represent temporal continuity, while similarity-based connections can capture relationships between acoustically related segments. This motivates the use of GraphSAGE to aggregate information across the segments of each track.

2.3. Transformer-Based Text Representation

Transformer architectures use self-attention to model relationships between elements in a sequence and have become widely used for contextual language representation (Vaswani et al., 2017). BERT introduced bidirectional Transformer-based pre-training for learning contextual representations that can be adapted to downstream language tasks (Devlin et al., 2019). DistilBERT is a smaller and computationally efficient BERT-family model obtained through knowledge distillation (Sanh et al., 2019).

In this project, DistilBERT is used as the BERT-based text encoder. Track titles are tokenized and passed through the frozen encoder to obtain contextual semantic representations. These representations provide complementary information to the acoustic features extracted from the audio graphs.

2.4. Multimodal Representation Fusion

Multimodal learning aims to combine information from different modalities so that complementary characteristics can contribute to a common prediction task. In the context of music understanding, acoustic information can describe properties such as timbre, rhythm, and spectral structure, while text can provide semantic information associated with the track.

A straightforward approach is to concatenate representations from different modalities before classification. More expressive fusion mechanisms can instead allow one modality to condition its representation on information from another modality. Cross-attention provides such a mechanism by allowing a query representation from one modality to attend to feature representations from another modality. In this work, both simple feature fusion and cross-attention are investigated to determine whether explicit interaction between textual and acoustic representations improves genre classification.

3. Dataset and Preprocessing

3.1. Dataset

This project uses the Free Music Archive (FMA) dataset, specifically the FMA-small subset, which contains 8,000 audio tracks together with genre metadata (Defferrard et al., 2017). The FMA-small subset is selected because it provides a manageable collection of publicly available music suitable for reproducible supervised learning experiments.

The classification task uses eight genre categories: Electronic, Experimental, Folk, Hip-Hop, Instrumental, International, Pop, and Rock. The original FMA metadata provides the track identifiers, genre labels, and track titles used in the experiments. The official dataset split is preserved throughout the experiments to avoid introducing additional train-test contamination.

The final dataset contains 6,400 training tracks, 800 validation tracks, and 800 test tracks. The official split is balanced across the eight selected genres, with 800 training tracks, 100 validation tracks, and 100 test tracks per genre before graph-generation filtering.

3.2. Task Definition

The original project specification considers multimodal music understanding using both audio and textual information. In the implemented experiment, the primary supervised task is formulated as an eight-class music genre classification problem.

For each track, the input consists of an audio representation and a textual representation derived from the track title. The target variable is the corresponding genre label. Formally, for a track i , the input can be represented as

$$T_i = (X_i^{audio}, X_i^{text}, y_i), \quad (1)$$

where X_i^{audio} denotes the acoustic representation, X_i^{text} denotes the title representation, and $y_i \in \{0, \dots, 7\}$ denotes the genre label.

Although the broader project specification allows multi-label tagging and other music-context prediction tasks, the implemented experiment uses single-label eight-class genre classification because the available metadata consistently provides a primary genre label and track title for the selected subset.

3.3. Audio Preprocessing

Each audio track is resampled to 22,050 Hz before feature extraction. The audio signal is then divided into fixed-length segments of 5 seconds. Only complete 5-second segments are retained so that every graph node corresponds to the same temporal duration.

For each segment, a 128-bin log-Mel spectrogram is computed. The resulting time-frequency representation is summarized along the temporal dimension to obtain a fixed-dimensional feature vector of size 128. Thus, if a track contains N complete five-second segments, its acoustic representation is expressed as

$$X^{audio} = [x_1, x_2, \dots, x_N], \quad x_i \in \mathbb{R}^{128}. \quad (2)$$

This representation provides a compact description of the spectral characteristics of each local audio segment while preserving the segment-level structure required for graph construction.

3.4. Audio Graph Construction

Each track is represented as a graph $G = (V, E)$, where each node corresponds to one complete five-second audio segment. The node feature matrix is therefore

$$X_G \in \mathbb{R}^{N \times 128}, \quad (3)$$

where N is the number of valid audio segments in the track.

Two types of edges are used to capture complementary relationships between segments. First, bidirectional temporal edges are added between consecutive segments to represent the sequential structure of the audio. Second, cosine-similarity edges are added between acoustically related segments. For each node, the two most similar other segments are selected according to cosine similarity.

The combination of temporal and similarity-based edges allows the graph to represent both local sequential continuity and non-local acoustic relationships. Consequently, the graph structure is not restricted to a simple chain and can connect segments that exhibit similar spectral characteristics even when they occur at different positions within the track.

Duplicate edges are removed during graph construction to maintain a deterministic graph representation. The resulting graph is stored using PyTorch Geometric’s graph data format, with node features, edge indices, genre labels, and track identifiers preserved for subsequent training and evaluation.

3.5. Text Preprocessing

The textual modality is represented using the track title associated with each FMA track. Titles are tokenized using the DistilBERT tokenizer with the pretrained model `distilbert-base-uncased` (Sanh et al., 2019).

A maximum sequence length of 32 tokens is used because the track titles in this dataset are substantially shorter than typical document-level inputs. Titles longer than the specified limit are truncated, while shorter titles are padded according to the tokenizer requirements.

The DistilBERT encoder is kept frozen during training. Its contextual representation is extracted from the encoded title and subsequently passed to a trainable classification or fusion component. Freezing the encoder reduces computational requirements and ensures that the experiment primarily evaluates the contribution of the learned multimodal fusion architecture.

3.6. Feature Normalization

To avoid information leakage between the dataset splits, feature normalization is fitted using only the training graphs. A `StandardScaler` is incrementally fitted over the node features of the training graphs and then applied without refitting to the validation and test graphs.

For an individual feature dimension, the normalized value is computed as

$$\hat{x} = \frac{x - \mu_{train}}{\sigma_{train}}, \quad (4)$$

where μ_{train} and σ_{train} are estimated exclusively from the training data.

This procedure ensures that information from the validation and test sets is not used during preprocessing. The fitted scaler is stored and reused consistently across the GraphSAGE and multimodal experiments.

3.7. Graph Generation and Data Availability

The official train, validation, and test split assignments are retained throughout preprocessing. During graph generation, six training tracks were excluded because their audio files were missing, unreadable, or did not contain a complete five-second segment under the fixed-segment preprocessing procedure. The excluded track identifiers are 98565, 98567, 98569, 99134, 108925, and 133297.

After this filtering step, the final graph dataset contains 6,394 training graphs, 800 validation graphs, and 800 test graphs. The validation and test sets remain aligned with the official split, while the reduction in the training set is explicitly accounted for during multimodal representation alignment.

For reproducibility, representative graph files are included in the project repository, while the complete processed graph collection and generated textual embeddings are treated as derived project artifacts.

4. Methodology

4.1. Overall Architecture

The proposed system combines graph-based acoustic representation learning with Transformer-based textual representation learning for music genre classification. Each music track is represented through two complementary modalities: an audio graph constructed from fixed-length audio segments and a textual representation obtained from the track title.

The overall architecture consists of three main components. First, GraphSAGE is used to encode the audio graph and

produce a track-level acoustic representation. Second, a frozen DistilBERT encoder produces a contextual representation of the track title. Finally, the two modality-specific representations are combined using either simple feature fusion or a cross-attention mechanism before classification.

For a track i , the two modality-specific representations can be written as

$$h_i^{audio} = \text{GNN}(G_i, X_i), \quad (5)$$

and

$$h_i^{text} = \text{DistilBERT}(X_i^{text}), \quad (6)$$

where G_i and X_i denote the audio graph and its node features, respectively. The resulting representations are then combined to obtain a multimodal representation used for genre prediction.

4.2. GraphSAGE Audio Encoder

The acoustic representation is learned using GraphSAGE (Hamilton et al., 2017). GraphSAGE performs neighborhood aggregation so that each node can incorporate information from its connected neighbors. This is particularly suitable for the constructed music graphs because the edges represent both temporal continuity and acoustic similarity between audio segments.

Let $h_v^{(l)}$ denote the representation of node v at layer l . A GraphSAGE layer aggregates the representations of neighboring nodes and combines the aggregated information with the node’s current representation. The general update can be expressed as

$$h_v^{(l+1)} = \sigma \left(W^{(l)} \left[h_v^{(l)} \parallel \text{AGG} \left(\{h_u^{(l)} : u \in \mathcal{N}(v)\} \right) \right] \right), \quad (7)$$

where $\mathcal{N}(v)$ represents the neighborhood of node v , AGG is a neighborhood aggregation operation, $\parallel$ denotes concatenation, $W^{(l)}$ is a learnable transformation, and σ is a nonlinear activation function.

The implemented GraphSAGE encoder contains two GraphSAGE convolution layers. The first layer transforms the 128-dimensional input node features into a 128-dimensional hidden representation. ReLU activation and dropout are applied after the first layer. A second GraphSAGE layer further propagates information across the graph.

After message passing, node representations are aggregated using global mean pooling to obtain a fixed-dimensional representation for the complete track:

$$h_i^{audio} = \text{MeanPool} \left(\{h_v^{(2)} : v \in V_i\} \right). \quad (8)$$

The resulting 128-dimensional vector represents the acoustic characteristics of the complete track.

4.3. DistilBERT Text Encoder

The textual modality is encoded using the pretrained `distilbert-base-uncased` model (Sanh et al., 2019). For each track, the title is tokenized and passed through the Transformer encoder. The resulting contextual representation is used as the semantic representation of the track.

The DistilBERT encoder is frozen during training. Therefore, its parameters are not updated using the genre-classification objective. Instead, the extracted 768-dimensional representation is passed to a trainable projection or classification component.

Let $h_i^{BERT} \in \mathbb{R}^{768}$ denote the representation obtained from DistilBERT. For multimodal fusion, this representation is projected into the same 128-dimensional feature space used by the audio encoder:

$$h_i^{text} = W_t h_i^{BERT} + b_t, \quad (9)$$

where

$$W_t \in \mathbb{R}^{128 \times 768}. \quad (10)$$

This projection makes the textual representation compatible with the graph-based acoustic representation.

4.4. Simple GNN-BERT Fusion

As a baseline multimodal fusion strategy, the GraphSAGE and DistilBERT representations are concatenated directly. The fused representation is defined as

$$z_i = h_i^{audio} \parallel h_i^{text}. \quad (11)$$

Since both representations have dimensionality 128 after projection, the resulting fused representation has dimensionality 256.

The fused representation is passed through a fully connected classification network consisting of a 256-to-128 linear layer, ReLU activation, dropout, and a final linear layer producing eight genre logits. The predicted class is obtained from the maximum output logit.

This simple fusion model provides an important baseline for evaluating whether combining acoustic and textual in-

formation is beneficial even without explicitly modeling interactions between the two modalities.

4.5. Cross-Attention GNN-BERT Fusion

The main multimodal architecture extends simple feature concatenation with cross-attention. The motivation is to allow the textual representation to selectively attend to different audio segments rather than treating all graph-level acoustic information equally.

After the second GraphSAGE layer, the node representations are retained as a sequence of segment-level features. Because different tracks contain different numbers of five-second segments, the variable-length node sequences are converted into a padded dense representation with an associated node mask.

The projected textual representation is used as the query of a multi-head attention mechanism, while the graph node representations are used as both keys and values. Formally, the projected textual representation is

$$q_i = W_t h_i^{BERT}, \quad (12)$$

and the audio node representations form the key and value sequences

$$K_i = V_i = [h_{i,1}, h_{i,2}, \dots, h_{i,N_i}]. \quad (13)$$

The cross-attention operation can then be expressed conceptually as

$$\text{Attn}(q_i, K_i, V_i) = \text{softmax} \left(\frac{q_i K_i^T}{\sqrt{d}} \right) V_i, \quad (14)$$

where d denotes the attention feature dimension.

The implementation uses four attention heads. The resulting attended textual representation is combined with the original projected text representation through a residual connection and layer normalization:

$$\tilde{h}_i^{text} = \text{LayerNorm} \left(\text{CrossAttn}(q_i, K_i, V_i) + h_i^{text} \right). \quad (15)$$

In parallel, a global mean pooling operation is applied to the GraphSAGE node representations to obtain a global acoustic representation. The final multimodal representation is then constructed as

$$z_i = \tilde{h}_i^{text} \parallel h_i^{audio}. \quad (16)$$

The resulting 256-dimensional representation is passed through a fully connected classification head to predict the eight genre classes.

4.6. Classification Objective

All classification models are trained using the standard multiclass cross-entropy loss. Given the predicted logits $\hat{y}_i$ and the ground-truth class y_i , the loss is defined as

$$\mathcal{L}_{CE} = -\frac{1}{N} \sum_{i=1}^N \log \left(\frac{\exp(\hat{y}_{i,y_i})}{\sum_{c=1}^C \exp(\hat{y}_{i,c})} \right), \quad (17)$$

where N is the number of training examples and $C = 8$ is the number of genre classes.

The same primary classification objective is used for the CNN, GraphSAGE, simple GNN-BERT, and cross-attention GNN-BERT models. The DistilBERT baseline also uses the same genre classification objective, but only its trainable classification head is optimized while the pre-trained DistilBERT encoder remains frozen.

4.7. Model Variants

To evaluate the contribution of each component, five model configurations are compared:

1. **DistilBERT Title:** A frozen DistilBERT encoder followed by a trainable classification head using only track titles.
2. **CNN:** A convolutional neural network operating directly on the log-Mel spectrogram representation of the first complete five-second audio segment.
3. **GraphSAGE:** A graph-based audio model operating on all complete five-second segments and their temporal and similarity-based graph connections.
4. **Simple GNN-BERT:** Concatenation-based fusion of the GraphSAGE audio representation and the projected DistilBERT text representation.
5. **Cross-Attention GNN-BERT:** The proposed multimodal model in which the textual representation attends to segment-level GraphSAGE representations before classification.

This progression provides a controlled comparison from unimodal text and audio baselines to graph-based audio modeling and finally multimodal fusion. It also allows the effect of simple fusion to be separated from the additional contribution of cross-attention.

5. Experimental Setup

5.1. Data Splits

All experiments use the official FMA-small train, validation, and test partitions. The original split contains 6,400 training tracks, 800 validation tracks, and 800 test tracks. The same partitions are maintained across the text, CNN, GraphSAGE, and multimodal experiments to ensure that model comparisons are performed on consistent evaluation sets.

During graph construction, six training tracks were excluded because their audio files were missing, unreadable, or did not contain a complete five-second segment under the fixed-segment preprocessing procedure. Therefore, the final graph-based training set contains 6,394 graphs, while the validation and test sets contain 800 graphs each.

The exclusion is restricted to the training set and does not modify the official validation or test partitions. Track identifiers are retained throughout preprocessing so that graph representations and textual embeddings can be aligned correctly for multimodal training.

5.2. Training Configuration

All neural network models are trained using the AdamW optimizer. Unless otherwise specified, a learning rate of 10^{-3} and weight decay of 10^{-4} are used. Dropout with a probability of 0.3 is applied in the trainable neural network components to reduce overfitting.

The GraphSAGE model is trained for 30 epochs using a cross-entropy classification objective. Model selection is performed using validation accuracy, and the checkpoint corresponding to the best validation performance is retained for final test evaluation.

The CNN baseline is trained for 8 epochs using the same optimizer and learning-rate configuration. The CNN operates on the first complete five-second audio segment available for each track.

The DistilBERT text baseline is trained for 20 epochs. The pretrained DistilBERT encoder is frozen, and only the trainable classification head is optimized. This reduces the number of trainable parameters and makes the experiment computationally feasible on the available CPU-based environment.

Both multimodal models are trained for 20 epochs. The simple GNN-BERT fusion model and the cross-attention GNN-BERT model use the same basic optimizer configuration so that differences in performance can be attributed primarily to the fusion strategy.

5.3. Model Selection

Validation performance is used to select the best checkpoint for each model. The model state corresponding to the highest validation accuracy is saved and subsequently evaluated once on the held-out test set.

This procedure prevents the test set from influencing model selection. The test set is used only after training and checkpoint selection have been completed.

For the cross-attention GNN-BERT model, the best validation performance is obtained at epoch 8. The corresponding checkpoint is used for the final evaluation. The best simple GNN-BERT checkpoint is obtained at epoch 16, while the best GraphSAGE checkpoint is obtained at epoch 11.

5.4. Evaluation Metrics

Because the task is an eight-class single-label classification problem, several complementary metrics are reported.

Accuracy. Accuracy measures the proportion of test examples for which the predicted genre is identical to the ground-truth genre:

$$\text{Accuracy} = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(\hat{y}_i = y_i), \quad (18)$$

where $\mathbb{I}(\cdot)$ is an indicator function.

Macro-F1. Macro-F1 calculates the F1 score independently for each class and then averages the resulting scores:

$$\text{Macro-F1} = \frac{1}{C} \sum_{c=1}^C F1_c. \quad (19)$$

This metric gives equal importance to every genre and is therefore useful for identifying whether a model performs well only on a small number of classes.

Micro-F1. Micro-F1 aggregates the contributions of all classes before calculating the F1 score. For the balanced single-label classification setting used in this project, Micro-F1 is numerically equivalent to accuracy.

Macro AUC-PR. Macro AUC-PR measures the area under the precision-recall curve separately for each class and then averages the class-level scores. It provides an additional evaluation of the quality of the predicted probability distributions, particularly when class-level prediction confidence differs across genres.

5.5. Baseline Comparisons

Five models are evaluated to understand the contribution of the different modalities and architectural components:

1. **DistilBERT Title Baseline:** uses only the track title as textual input.
2. **CNN Audio Baseline:** uses a convolutional network operating on the first complete five-second log-Mel segment.
3. **GraphSAGE Audio Model:** uses all available complete five-second segments represented as nodes in the constructed audio graph.
4. **Simple GNN-BERT Fusion:** combines the GraphSAGE and DistilBERT representations through direct concatenation.
5. **Cross-Attention GNN-BERT:** combines the two modalities using multi-head cross-attention between the textual representation and segment-level graph representations.

The baseline progression is designed to answer three questions. First, how informative is the textual modality by itself? Second, does graph-based audio modeling provide an advantage over a conventional CNN baseline? Third, does combining acoustic and textual representations improve performance beyond either modality individually, and does cross-attention provide an additional benefit over simple concatenation?

5.6. Implementation Details

The experiments are implemented using PyTorch and PyTorch Geometric. Graph construction and graph-level batching are handled using PyTorch Geometric, while the pre-trained DistilBERT tokenizer and encoder are provided by the Transformers library.

The audio processing pipeline uses a sampling rate of 22,050 Hz and 128-bin log-Mel features. Graph nodes correspond to complete five-second segments, and graph edges combine bidirectional temporal connections with the top-two cosine-similarity relationships for each node.

The GraphSAGE hidden dimension is 128. The multimodal models project the 768-dimensional DistilBERT representation into the same 128-dimensional space. The cross-attention module uses four attention heads. The final multimodal representation has 256 dimensions after concatenating the attended textual representation with the global graph representation.

All experiments are conducted using a CPU-based training environment. To improve reproducibility, the preprocess-

Table 1. Test-set performance of the evaluated models. Best results are shown in bold.

Model	Acc.	Macro-F1	Micro-F1	AUC-PR
DistilBERT	26.62	24.99	26.62	25.51
CNN	31.50	29.40	31.50	32.48
GraphSAGE	37.13	35.88	37.13	38.72
Simple GNN-BERT	40.38	40.78	40.38	42.71
Cross-Attn. GNN-BERT	40.50	41.18	40.50	43.25

ing scaler, model checkpoints, representative graph files, generated embeddings, and evaluation outputs are saved as project artifacts.

6. Results

6.1. Overall Model Comparison

Table 1 summarizes the performance of all five evaluated models on the held-out test set. The results show a clear progression from unimodal baselines toward multimodal architectures.

The frozen DistilBERT title model achieves 26.62% accuracy and 24.99% Macro-F1, indicating that track titles alone provide useful but limited information for predicting musical genre. The CNN baseline performs better, reaching 31.50% accuracy and 29.40% Macro-F1. This indicates that acoustic information contained in the log-Mel representation provides stronger genre-related information than the title-only textual representation in this experimental setting.

The GraphSAGE model further improves performance to 37.13% accuracy and 35.88% Macro-F1. Its improvement over the CNN baseline suggests that representing the audio as a graph can capture relationships between different temporal segments that are not explicitly modeled by the conventional convolutional baseline. However, this comparison should be interpreted with care because the two models use different amounts of audio context: the CNN baseline uses only the first complete five-second segment, whereas GraphSAGE processes all complete five-second segments available for each track.

The multimodal models achieve the strongest overall performance. Simple GNN-BERT fusion increases accuracy to 40.38% and Macro-F1 to 40.78%. The cross-attention GNN-BERT model achieves the best overall result, with 40.50% accuracy, 41.18% Macro-F1, 40.50% Micro-F1, and 43.25% Macro AUC-PR.

These results support the main hypothesis that combining acoustic and textual representations is more effective than using either modality independently. The improvement from GraphSAGE to simple GNN-BERT is particularly notable,




Figure 1. Overall test-set performance comparison of the evaluated DistilBERT, CNN, GraphSAGE, Simple GNN-BERT, and Cross-Attention GNN-BERT models.

Table 2. Ablation of multimodal fusion strategies.

Fusion Method	Accuracy	Macro-F1	Micro-F1	AUC-PR
Simple Fusion	40.38	40.78	40.38	42.71
Cross-Attention	40.50	41.18	40.50	43.25

indicating that textual information provides complementary information beyond the graph-based acoustic representation

6.2. Fusion Ablation

To examine whether explicit cross-modal interaction provides an advantage over direct concatenation, the two multimodal fusion strategies are compared in Table 2.

Cross-attention provides a modest improvement over simple concatenation. Accuracy increases by 0.12 percentage points, while Macro-F1 increases by approximately 0.41 percentage points. Macro AUC-PR increases by approximately 0.54 percentage points.

The relatively small difference between the two fusion mechanisms suggests that the primary benefit comes from combining the two modalities rather than from the specific attention mechanism itself. Nevertheless, cross-attention achieves the best value across all reported metrics, indicating that allowing the textual representation to selectively interact with segment-level acoustic features provides a useful, although limited, additional benefit.

6.3. Class-Level Performance

The class-level results reveal substantial variation in genre recognition performance. The Cross-Attention GNN-BERT model achieves its strongest F1 scores on Hip-Hop (64.71%), International (57.75%), and Rock (53.76%).

These classes appear to contain acoustic and/or semantic characteristics that are comparatively easier for the multimodal representation to distinguish.

In contrast, Folk (21.77%), Pop (21.05%), and Experimental (26.34%) remain more challenging. These classes may contain substantial intra-class variation or acoustic similarities with other genres. In particular, the relatively low performance on Folk and Pop indicates that genre boundaries are not always clearly represented by the available audio and title information.

The GraphSAGE model also performs particularly well on Hip-Hop and Rock, with F1 scores of 60.87% and 55.77%, respectively. This suggests that acoustic structure alone can provide strong evidence for some genres. However, multimodal fusion improves several classes where the unimodal models are less reliable.

6.4. Confusion Analysis

The confusion matrices show that genre difficulty varies across classes. GraphSAGE correctly classifies 70 Hip-Hop and 58 Rock test examples, while substantial confusion remains among Folk, Experimental, Instrumental, and International. The cross-attention model correctly identifies 60 Electronic, 27 Experimental, 27 Folk, 55 Hip-Hop, 33 Instrumental, 54 International, 18 Pop, and 50 Rock examples. Folk is frequently confused with Experimental and Instrumental, while Pop receives predictions from several genres, indicating overlapping acoustic and semantic characteristics.

6.5. Main Finding

Overall, multimodal fusion improves music genre classification. The cross-attention GNN-BERT model achieves the best performance, improving accuracy from 26.62

The improvement over simple fusion is comparatively small, indicating that most of the gain comes from combining audio and text, while cross-attention provides an additional but modest benefit.

7. Representation Analysis

7.1. t-SNE Visualization

To investigate how the learned representations organize the different music genres, t-SNE visualizations are generated for the audio, textual, and fused representations of the 800 test tracks (van der Maaten & Hinton, 2008). t-SNE is applied independently to each representation using a two-dimensional embedding with perplexity 30, PCA initialization, automatic learning rate selection, and random seed 42.

Figure 2 compares the resulting two-dimensional representations. The audio representation corresponds to the 128-dimensional GraphSAGE track-level embedding, while the textual representation corresponds to the projected 128-dimensional DistilBERT embedding. The fused representation corresponds to the 256-dimensional representation obtained by combining the attended textual representation with the global acoustic representation of the cross-attention GNN-BERT model.

The audio representation exhibits stronger organization for several acoustically distinctive genres, particularly Hip-Hop and Rock. This observation is consistent with the relatively high class-level F1 scores achieved by GraphSAGE for these genres. However, substantial overlap remains among genres such as Folk, Experimental, Instrumental, and International, indicating that their acoustic characteristics are not completely separable in the learned embedding space.

The textual representation shows a different organization from the audio representation. Since the text modality is based only on track titles, its semantic information is relatively limited compared with lyric- or caption-based representations. Nevertheless, some genre-related structure is present, suggesting that titles can provide useful semantic cues even without additional textual context.

The fused representation provides a more informative joint representation by combining the complementary information from both modalities. Several genre groups become more structured after fusion, although considerable overlap remains. This behavior is consistent with the quantitative results: the multimodal models outperform both the text-only and audio-only baselines, but the classification problem remains challenging for several genre categories.

The visualization therefore supports the quantitative finding that the two modalities capture complementary information. At the same time, the remaining overlap demonstrates that multimodal fusion does not produce perfectly separable genre representations and that additional information, such as lyrics, richer musical structure, or larger-scale audio representations, could potentially improve performance.

8. Case Studies

Three test-set examples are examined to illustrate how the acoustic and textual modalities contribute to multimodal genre prediction. The cases represent complementary fusion behavior: improvement over both unimodal models, strong acoustic evidence, and correction of an acoustic-only error.

8.1. Case Study 1: Helicopter de Cristal

Track 1066, “Helicopter de Cristal,” has Experimental as the ground-truth genre. DistilBERT predicts International

(39.48

This example shows that combining the two modalities can recover the correct genre when neither unimodal model provides the correct top prediction.

8.2. Case Study 2: New Wave Hippies

Track 4108, “New Wave Hippies,” belongs to Rock. DistilBERT predicts Experimental with 44.43

This case demonstrates that the acoustic graph can provide strong genre evidence, which is preserved by multimodal fusion. The result also shows that textual information does not always need to dominate when the acoustic representation is highly informative.

8.3. Case Study 3: Mano De Dios

Track 704, “Mano De Dios,” has International as the ground truth. DistilBERT correctly predicts International with 44.51

This example highlights how textual information can complement the acoustic graph and help correct an acoustic-only error.

8.4. Case Study Summary

The three cases illustrate different benefits of multimodal fusion. In the first, fusion corrects errors from both unimodal models; in the second, it preserves a strong acoustic signal; and in the third, textual information helps correct an acoustic prediction.

Overall, the examples support the view that audio and text provide complementary information whose contribution varies across tracks. However, fusion does not guarantee correct predictions, as reflected by the overall test performance.

9. Discussion

9.1. Effect of Multimodal Fusion

Combining acoustic and textual information improves performance over both unimodal models. DistilBERT achieves 26.62

This improvement indicates that the two modalities provide complementary information. Audio graphs capture spectral and temporal characteristics, while track titles provide semantic cues. The case studies further show that fusion can correct errors made by individual modalities, demonstrating the practical benefit of combining both sources.




Figure 2. t-SNE visualization of the test-set representations. The audio representation is obtained from the GraphSAGE encoder, the text representation is obtained from the projected DistilBERT embedding, and the fused representation combines acoustic and textual information through cross-attention.

9.2. Contribution of Cross-Attention

Cross-attention achieves the best overall performance, but its improvement over simple fusion is modest. Accuracy increases from 40.38

This suggests that most of the performance gain comes from multimodal combination rather than the attention mechanism itself. The limited textual information, consisting only of track titles, may also restrict the benefit of cross-attention. Therefore, cross-attention is useful, but it is not substantially superior to simple fusion under the current experimental setting.

9.3. Importance of Graph-Based Audio Representation

GraphSAGE outperforms the CNN baseline, achieving 37.13

However, the comparison is not fully controlled. CNN uses only the first complete five-second segment, whereas GraphSAGE processes all complete segments available for each track. Therefore, the improvement cannot be attributed entirely to graph-based message passing. A future controlled comparison should provide both models with the same audio context.

9.4. Class-Level Behavior

The cross-attention model performs particularly well on Hip-Hop, International, and Rock, with F1 scores of 64.71

The confusion patterns show overlap among several genres, particularly Folk, Experimental, and Instrumental. These results also highlight the importance of Macro-F1 alongside accuracy, since accuracy alone may hide substantial differences between individual classes.

9.5. Limitations

The study has several limitations. First, the textual modality is restricted to track titles rather than richer sources such as lyrics, tags, or captions. Second, the DistilBERT encoder is frozen, preventing task-specific adaptation. Third, CNN and GraphSAGE receive different amounts of audio context, limiting the fairness of their comparison.

In addition, six training tracks were excluded during graph construction because their audio files were missing, unreadable, or too short for a complete five-second segment. Finally, the experiment focuses on eight-class genre classification and does not evaluate the broader multi-label tagging, emotion regression, or cross-modal retrieval objectives.

9.6. Implications and Future Work

Future work could incorporate richer textual information such as lyrics, tags, and MusicCaps-style descriptions, while fine-tuning the language encoder for the music domain. The audio representation could also be improved using beat-synchronous segmentation, chroma features, learned spectrogram encoders, or longer temporal contexts.

A controlled CNN–GraphSAGE comparison using identical audio segments would better isolate the contribution of graph-based modeling. The framework could also be extended to multi-label music tagging, emotion prediction, and audio–text retrieval to evaluate music-context understanding beyond single-label genre classification.

10. Conclusion

This project investigated multimodal music genre classification by combining graph-based acoustic representations with DistilBERT-based textual representations. Audio tracks were represented as graphs of five-second log-Mel segments, while track titles were encoded using a frozen Distil-

BERT model. Among the evaluated approaches, the Cross-Attention GNN-BERT model achieved the best test performance, with 40.50% accuracy, 41.18% Macro-F1, and 43.25% Macro AUC-PR. The results show that combining acoustic and textual information provides a larger improvement than either modality alone, while cross-attention offers a smaller additional gain over simple fusion. Future work could extend the framework to multi-label music tagging, emotion prediction, richer textual sources such as lyrics or captions, and audio-text retrieval.

References

- Defferrard, M., Benzi, K., Vandergheynst, P., and Bresson, X. Fma: A dataset for music analysis. In *Proceedings of the 18th International Society for Music Information Retrieval Conference (ISMIR)*, pp. 316–323, Suzhou, China, 2017.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. BERT: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies*, volume 1, pp. 4171–4186, Minneapolis, Minnesota, 2019. Association for Computational Linguistics. doi: 10.18653/v1/N19-1423.
- Hamilton, W. L., Ying, Z., and Leskovec, J. Inductive representation learning on large graphs. In *Advances in Neural Information Processing Systems*, volume 30, 2017.
- Sanh, V., Debut, L., Chaumond, J., and Wolf, T. Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter. In *NeurIPS EMC² Workshop*, 2019.
- van der Maaten, L. and Hinton, G. Visualizing data using t-sne. *Journal of Machine Learning Research*, 9:2579–2605, 2008.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. Attention is all you need. In *Advances in Neural Information Processing Systems*, volume 30, 2017.